

Parallel Programming in OpenMP – part II

Outline

- ❑ Data scoping – cont'd
- ❑ Orphaning
- ❑ Tasking
- ❑ OpenMP correctness & Data Races
- ❑ Runtime library
- ❑ Scheduling
- ❑ A real world example

OpenMP Syntax

More on data scoping

OpenMP Syntax

Reminder: the “private” clause –

- ❑ declares variables private to each thread:

```
#pragma omp directive private (list)
```

- ❑ i.e. a **new** variable is declared once for each thread
- ❑ all references are replaced with references to the newly declared variable
- ❑ variables declared private are uninitialized for each thread!

OpenMP Syntax

Consequences of private(...):

```
main() {  
    ...  
    A = 10;  
  
    #pragma omp parallel  
    {  
        #pragma omp for private(i, A, B) ...  
        for(i = 0; i < n; i++) {  
            ...  
            B = A + i;    // A undefined!  
                        // unless declared firstprivate  
  
            ...  
        } /* end of omp for */  
  
        ...  
    } /* end of omp parallel */  
  
    C = B;                // B undefined!  
                        // unless declared lastprivate  
  
}
```

OpenMP Syntax

Solutions:

```
#pragma omp ... firstprivate(list)
```

- ❑ All variables in list are initialized with the value the original object had before entering the parallel construct.

```
#pragma omp ... lastprivate(list)
```

- ❑ The thread that executes the sequentially last iteration updates all variables in list.

OpenMP Syntax

The “threadprivate” and “copyin” clauses:

- ❑ **threadprivate(list)**: creates a private copy of global data (e.g. common blocks or global variables in modules in Fortran) for each thread
- ❑ **copyin(list)**: copies the values from the master thread into the private copies – like a ‘firstprivate’ for global variables
- ❑ subsequent modifications of list affect only the private copies – within one parallel region

OpenMP Syntax

Example 1:

```
int counter = 0;
#pragma omp threadprivate(counter)

int increment_counter()
{
    counter++;
    return(counter);
}
```

```
INTEGER FUNCTION INCREMENT_COUNTER()
COMMON/A22_COMMON/COUNTER
!$OMP THREADPRIVATE (/A22_COMMON/)
COUNTER = COUNTER +1
INCREMENT_COUNTER = COUNTER
RETURN
END FUNCTION INCREMENT_COUNTER
```


OpenMP Syntax

Example 2:

```
int
increment_counter()
{
    static int counter = 0;
    #pragma omp threadprivate(counter)

    counter++;
    return(counter);
}
```

OpenMP Syntax

The `copyprivate(...)` clause

- ❑ copying a value out of a single region into the private data of other threads (“broadcast”)

```
#pragma omp single copyprivate(list)
{
    ...
}
```

```
!$OMP SINGLE ....
```

```
...
!$OMP END SINGLE COPYPRIVATE (LIST)
```

OpenMP Syntax

Example:

```
int x, y; /* global data */
#pragma omp threadprivate(x, y)

void use_values(int id, int a, int b) {
    printf("  TID %d: a = %d, b = %d, c = %d, d = %d\n",
          id, a, b, x, y);
}

void init(int id, int *a, int *b) {
    int r_a, r_b;

    #pragma omp single copyprivate(r_a, r_b, x, y)
    {
        scanf("%d %d %d %d", &r_a, &r_b, &x, &y);

        *a = r_a; *b = r_b;
        use_values(id, *a, *b);
    }

    ...
}
```

OpenMP Syntax

Example (cont'd):

```
int main(int argc, char *argv[] ) {  
  
    int tid = 0;  
    int a, b;  
  
    #pragma omp parallel private(tid,a,b)  
    {  
        #ifdef _OPENMP  
            tid = omp_get_thread_num();  
        #endif  
  
        init(tid, &a, &b);  
        printf("In main - TID %d: a = %d, b = %d, ",  
              " x = %d, y = %d\n",  
              tid, a, b, x, y);  
    } /* end of omp parallel */  
  
    return(0);  
}
```

OpenMP Syntax

Example output:

```
$ OMP_NUM_THREADS=3 ./copypriv
1 2 3 4
```

```
TID 2: a = 1, b = 2, c = 3, d = 4
```

```
TID 1: a = 1, b = 2, c = 3, d = 4
```

```
TID 0: a = 1, b = 2, c = 3, d = 4
```

```
In main - TID 0: a = 1, b = 2, x = 3, y = 4
```

```
In main - TID 1: a = 1, b = 2, x = 3, y = 4
```

```
In main - TID 2: a = 1, b = 2, x = 3, y = 4
```

```
env OMP_NUM_THREADS=3 ./copypriv
1 2 3 4
```

```
TID 1: a = 0, b = 1, c = 3, d = 4
```

```
TID 0: a = 0, b = 0, c = 3, d = 4
```

```
TID 2: a = 1, b = 2, c = 3, d = 4
```

```
In main - TID 2: a = 1, b = 2, x = 3, y = 4
```

```
In main - TID 0: a = 0, b = 0, x = 3, y = 4
```

```
In main - TID 1: a = 0, b = 1, x = 3, y = 4
```

without copyprivate
on r_a and r_b

OpenMP Orphaning

Orphaning in OpenMP

OpenMP Orphaning

The OpenMP standard does not restrict worksharing and synchronization directives to be within the lexical extent of a parallel region. Those directives can be orphaned, i.e. they can appear outside a parallel region:

```
#pragma omp parallel
{
    :
    dowork();
    :
}
```

orphaned
worksharing
directive

```
void dowork(void) {
    :
    #pragma omp for
    for(i=0; i<N; i++) {
        :
    }
}
```

OpenMP Orphaning

- ❑ When an orphaned directive is detected within the dynamic extent of a parallel region, its behaviour is similar to the non-orphaned case.
- ❑ When an orphaned directive is detected in the sequential part of the program, it will be ignored.

```
dowork(); // serial for  
  
#pragma omp parallel  
{  
    :  
    dowork(); // parallel for  
    :  
}
```

```
void dowork(void) {  
    :  
    #pragma omp for  
    for(i=0; i<N; i++) {  
        :  
    }  
}
```


OpenMP syntax

Tasking (since OpenMP 3.0)

- ❑ allows parallelization of work that is generated dynamically
- ❑ provides a flexible model for irregular parallelism
- ❑ uses a “task pool” concept
- ❑ new opportunities:
 - ❑ while loops
 - ❑ recursive structures

OpenMP syntax

❑ Syntax C/C++:

```
#pragma omp task [clause]
{
    ...
}
```

❑ clause can be

- ❑ if (int_expr)
- ❑ default(shared|none)
- ❑ private(list), shared(list)
- ❑ firstprivate(list)
- ❑ untied

OpenMP syntax

❑ Syntax Fortran:

```
!$OMP task [clause]  
  
...  
!$OMP end task
```

❑ where clause can be

- ❑ if (int_expr)
- ❑ default(shared|private|firstprivate|none)
- ❑ private(list), shared(list)
- ❑ firstprivate(list)
- ❑ untied

OpenMP syntax

Tasking example I:

while loop:

```
p = lhead;
while (p != NULL)
{
    do_work(p);
    p = next(p);
}
```

parallel while loop with OpenMP tasks:

```
#pragma omp parallel
{
    #pragma omp single
    {
        p = lhead;
        while (p != NULL) {
            #pragma omp task
            {
                do_work(p);
            }
            p = next(p);
        }
    } // end of single
} // end of parallel
```

OpenMP syntax

What's going on?

```
#pragma omp parallel
{
    #pragma omp single
    {
        p = lhead;
        while (p != NULL) {
            #pragma omp task
            {
                do_work(p);
            }
            p = next(p);
        }
    } // end of single
} // end of parallel
```

← start of parallel region

← one thread only, please


← task generation – tasks are added to the task list

← **all work is done here!**

← implicit barrier – all unfinished tasks have to be finished

OpenMP syntax

The task pool concept:



courtesy: Ruud van der Pas, Oracle

OpenMP syntax

Tasks and recursion: Fibonacci numbers

- ❑ Recursive scheme to calculate the n^{th} Fibonacci number:
 - ❑ $\text{fib}(n) = \text{fib}(n-1) + \text{fib}(n-2)$
 - ❑ stopping criterion: return 1 if $n < 2$
- ❑ Caveat: this method is not very effective, but used here to demonstrate the concept of tasking!

OpenMP syntax

The sequential code:

```
int
main(int argc, char* argv[]) {
    [...]
    fib(input);
    [...]
}
```

```
int
fib(int n) {

    int x, y;

    if (n < 2) return n;

    x = fib(n - 1);
    y = fib(n - 2);

    return(x + y);
}
```


OpenMP syntax

OpenMP version of fib() with tasks:

```
int
fib(int n) {

    int x, y;

    if (n < 2) return n;

    #pragma omp task shared(x)
    x = fib(n - 1);
    #pragma omp task shared(y)
    y = fib(n - 2);

    #pragma omp taskwait
    return(x + y);

}
```

note the special
scoping rules!

generate two tasks,
calling fib() recursively

task synchronization -
to get the right results

OpenMP syntax

Scoping rules with tasks:

- ❑ Static and global variables are shared
- ❑ Local (aka automatic) variables are private
- ❑ Orphaned task variables are firstprivate
- ❑ Non-orphaned task variables inherit the shared attribute
- ❑ (Local) Task variables are firstprivate, unless declared shared
- ❑ Thus, we have to declare x and y as shared

OpenMP syntax

Task synchronization:

- ❑ `#pragma omp taskwait`
- ❑ suspends the encountering task, until all child tasks are completed
- ❑ direct children only, not descendants
- ❑ needed here, to make sure that x and y are still exist when we take the sum.

OpenMP syntax

OpenMP version of main() with tasks:

```
int
main(int argc, char* argv[]) {
    [...]

    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    } // end of omp parallel


    [...]
}
```

start of parallel region -
team of worker threads

task generation by one
thread, only!

OpenMP syntax

Results of the Fibonacci program



courtesy: Chr. Terboven, RWTH Aachen

OpenMP syntax

Notes on the Fibonacci speedup results:

- ❑ The simple OpenMP version (omp-v1) doesn't scale – as expected – due to the large amount of tasks generated
- ❑ Improvement 1 (omp-v2):
 - ❑ add an if-clause to the tasks:
`#pragma omp task if(n>=30) shared(...)`
 - ❑ improves the speed-up, but still not perfect
- ❑ Improvement 2 (omp-v3): (see next slide)

OpenMP syntax

version omp-v3 of fib() with tasks:

```
int
fib(int n) {

    int x, y;

    if (n < 2) return n;
    if (n < 30) {
        return(fib(n-1) + fib(n-2));
    }

    #pragma omp task shared(x)
    x = fib(n - 1);
    #pragma omp task shared(y)
    y = fib(n - 2);

    #pragma omp taskwait
    return(x + y);
}
```

OpenMP syntax

Some notes on tasking:

- ❑ tasks allow to exploit parallelism with OpenMP, that hasn't been possible before
- ❑ this makes OpenMP a more powerful parallel programming API
 - ❑ goes beyond the scope in this course
- ❑ a nice intro about tasking in OpenMP:

<https://www.openmp.org/events/webinar-how-to-get-openmp-tasks-to-do-your-work/>

OpenMP: Error detection

Tools to check your OpenMP code

OpenMP compile-time checks

OpenMP checks supported by Oracle Studio compilers:

- ❑ -xloopinfo – info on loops
- ❑ -xvpara – compile time warnings on ...
 - ❑ ... scoping problems
 - ❑ ... possible data races

OpenMP: Data Race Detection

Oracle Studio has a data race detection tool:

- ❑ Thread Analyzer (`tha`)
- ❑ Quick usage guide:
 - ❑ compile and link with `-xinstrument=datarace`
 - ❑ run it: `collect -r on a.out`
 - ❑ view results (GUI): `tha tha.1.er`
 - ❑ or CLI: `er_print -races tha.1.er`

OpenMP: Data Race Detection

Example:

```
int main(int argc, char *argv[]) {  
  
    int i, total = 0, N = 2000000;  
    int primes[N];  
    #pragma omp parallel for  
    for( i = 2; i < N; i++ ) {  
        if ( is_prime(i) ) {  
            primes[total] = i;  
            total++;  
        }  
    }  
  
    printf("# of prime numbers between 2 and %d: %d\n",  
          N, total);  
    return(0);  
}
```

OpenMP: Data Race Detection

Example (cont'd): compile and run

```
$ suncc -g -fast -o prime prime.c
$ ptime ./prime
# of prime numbers between 2 and 2000000: 148933

real          10.862
user          10.483
sys           0.056

$ suncc -g -fast -xopenmp -xloopinfo -o prime prime.c
$ "prime.c", line 7: PARALLELIZED, user pragma used
$ OMP_NUM_THREADS=4 ptime ./prime
# of prime numbers between 2 and 2000000: 148310

real          3.662    <--- speed-up: 2.9x
user          10.494
sys           0.055
```

OpenMP: Data Race Detection

Example (cont'd): run, run, ... and use collect

```
$ OMP_NUM_THREADS=4 ./prime
# of prime numbers between 2 and 2000000: 148310

$ OMP_NUM_THREADS=4 ./prime
# of prime numbers between 2 and 2000000: 148328

$ suncc -g -fast -xopenmp -xinstrument=datarace \
  -o prime prime.c

$ OMP_NUM_THREADS=4 collect -r on ./prime
Creating experiment database tha.1.er ...
# of prime numbers between 2 and 2000000: 125581

$
```

OpenMP: Data Race Detection

Example (cont'd): analyze the collect data

```
$ er_print -races tha.1.er
Total Races:  2 Experiment:  tha.1.er

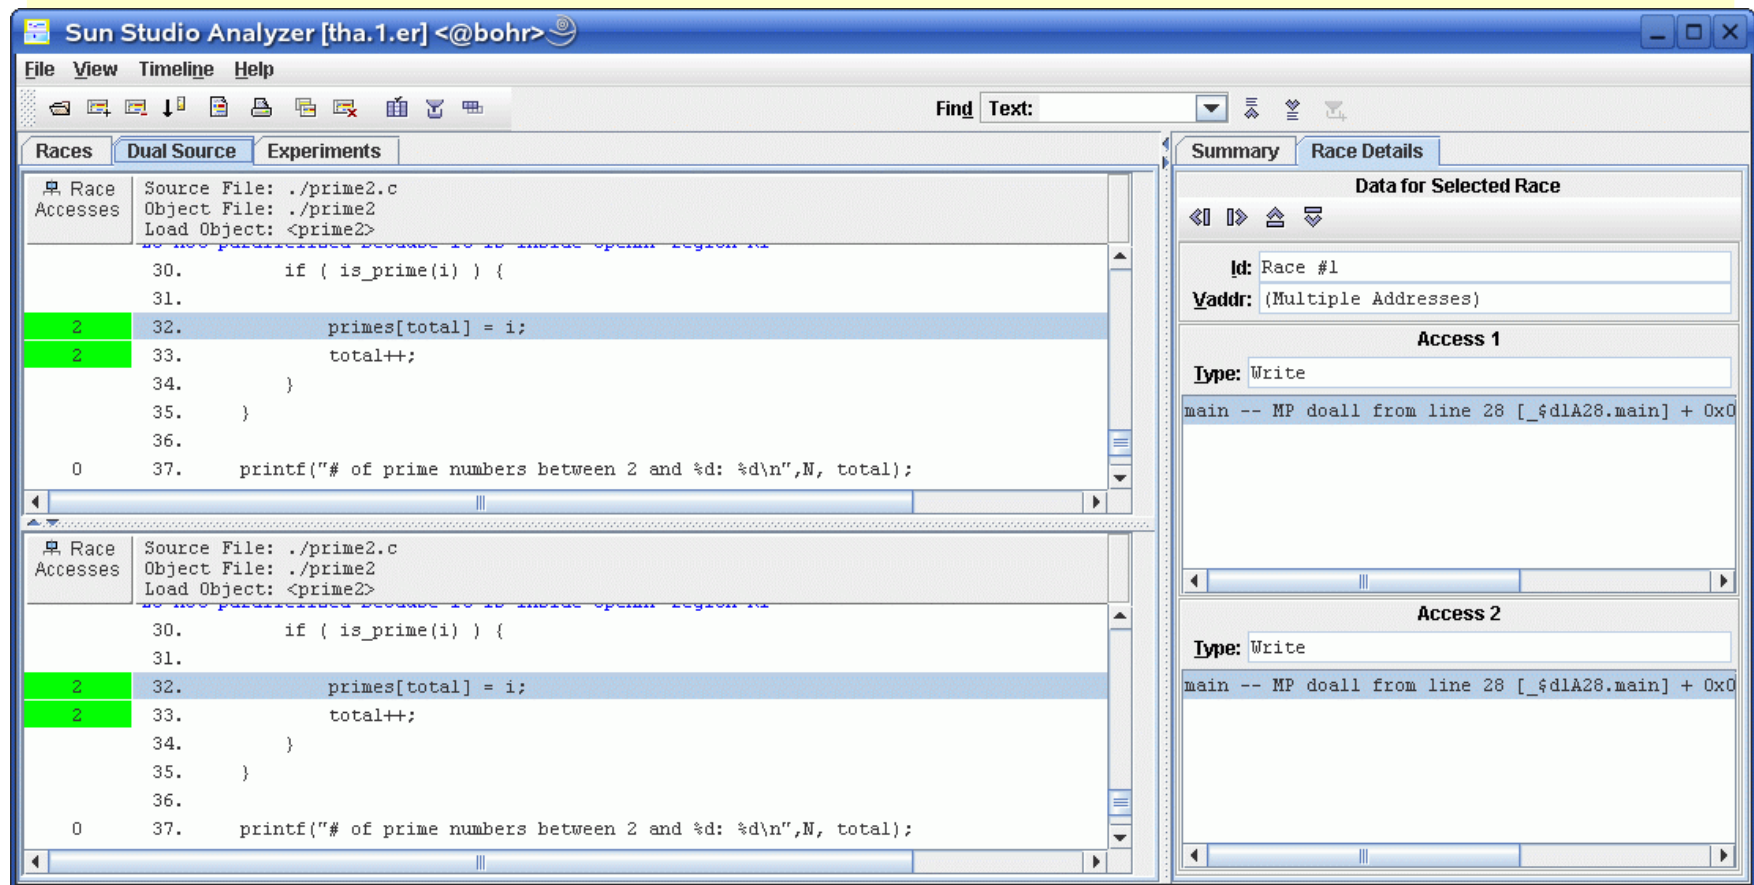
Race #1, Vaddr: (Multiple Addresses)
    Access 1: Write, main -- MP doall from line 6 ...
                line 10 in "prime.c"
    Access 2: Write, main -- MP doall from line 6 ...
                line 10 in "prime.c"
    Total Traces: 1

Race #2, Vaddr: 0xffbfecb4
    Access 1: Write, main -- MP doall from line 6 ...
                line 11 in "prime.c"
    Access 2: Write, main -- MP doall from line 6 ...
                line 11 in "prime.c"
    Total Traces: 1
```

OpenMP: Data Race Detection

Example (cont'd): analyze the collect data

```
$ tha tha.1.er
```



Sun Studio Analyzer [tha.1.er] <@bohr>

File View Timeline Help

Races Dual Source Experiments

Race Accesses

Source File: ./prime2.c
Object File: ./prime2
Load Object: <prime2>

```

30.         if ( is_prime(i) ) {
31.
2  32.         primes[total] = i;
2  33.         total++;
34.         }
35.     }
36.
0  37.     printf("# of prime numbers between 2 and %d: %d\n",N, total);
  
```

Summary Race Details

Data for Selected Race

Id: Race #1
Vaddr: (Multiple Addresses)

Access 1

Type: Write
main -- MP doall from line 28 [_\$dlA28.main] + 0x0

Access 2

Type: Write
main -- MP doall from line 28 [_\$dlA28.main] + 0x0

OpenMP: Data Race Detection

Example (cont'd): fix the bug

```
int main(int argc, char *argv[]) {  
  
    int i, total = 0, N = 2000000;  
    int primes[N];  
    #pragma omp parallel for  
  
    for( i = 2; i < N; i++ ) {  
        if ( is_prime(i) ) {  
            #pragma omp critical  
            { primes[total] = i;  
              total++;  
            }  
        }  
    }  
    printf("# of prime numbers between 2 and %d: %d\n",  
           N, total);  
    return(0);  
}
```

OpenMP: Data Race Detection

Example (cont'd): check – and recompile

```
$ suncc -g -fast -xopenmp -xinstrument=datarace \  
-o prime prime.c
```

```
$ OMP_NUM_THREADS=4 collect -r on ./prime  
Creating experiment database tha.2.er ...  
# of prime numbers between 2 and 2000000: 148933
```

```
$ er_print -races tha.2.er  
Total Races: 0 Experiment: tha.2.er
```

```
$ cc -g -fast -xopenmp -o prime prime.c  
$ OMP_NUM_THREADS=4 ptime ./prime  
# of prime numbers between 2 and 2000000: 148933
```

```
real      3.561  
user     10.393  
sys       0.051
```

OpenMP: Data Race Detection

- ❑ Some notes:
 - ❑ Oracle Studio does only support OpenMP up to version 4.0
 - ❑ most of the features we use here are supported, but not all
 - ❑ thus, using this tool will not work for everything
 - ❑ e.g. “doacross” loops

OpenMP Scheduling

Controlling the scheduling
of OpenMP threads

OpenMP Scheduling

Load balancing:

- ❑ Important aspect of performance
- ❑ Especially for less regular workloads, e.g.
 - ❑ transposing a matrix
 - ❑ multiplications of triangular matrices
 - ❑ parallel searches in a linked list
- ❑ The **schedule** clause provides different iteration scheduling algorithms for loops

OpenMP Scheduling

The “schedule” clause:

```
#pragma omp for schedule(static[,chunk])
```

```
#pragma omp for schedule(dynamic[,chunk])
```

```
#pragma omp for schedule(guided[,chunk])
```

```
#pragma omp for schedule(auto) - new in 3.0
```

```
#pragma omp for schedule(runtime)
```

- ❑ If there is no schedule clause, the default is static.

OpenMP Scheduling

```
#pragma omp for schedule(static[,chunk])
```

Static schedule:

- ❑ Iterations are divided into pieces of size chunk and then **statically** assigned to the threads.
- ❑ If chunk is not defined, the work (N) is equally divided among the number of threads (P), i.e. $\text{chunk} = N/P$.

OpenMP Scheduling

```
#pragma omp for schedule(dynamic[, chunk])
```

Dynamic schedule:

- ❑ Iterations are divided into pieces of size chunk and then **dynamically** assigned to the threads – i.e. when a thread has finished one chunk, it is assigned a new one.
- ❑ The default chunk size is 1.

OpenMP Scheduling

```
#pragma omp for schedule(guided[,chunk])
```

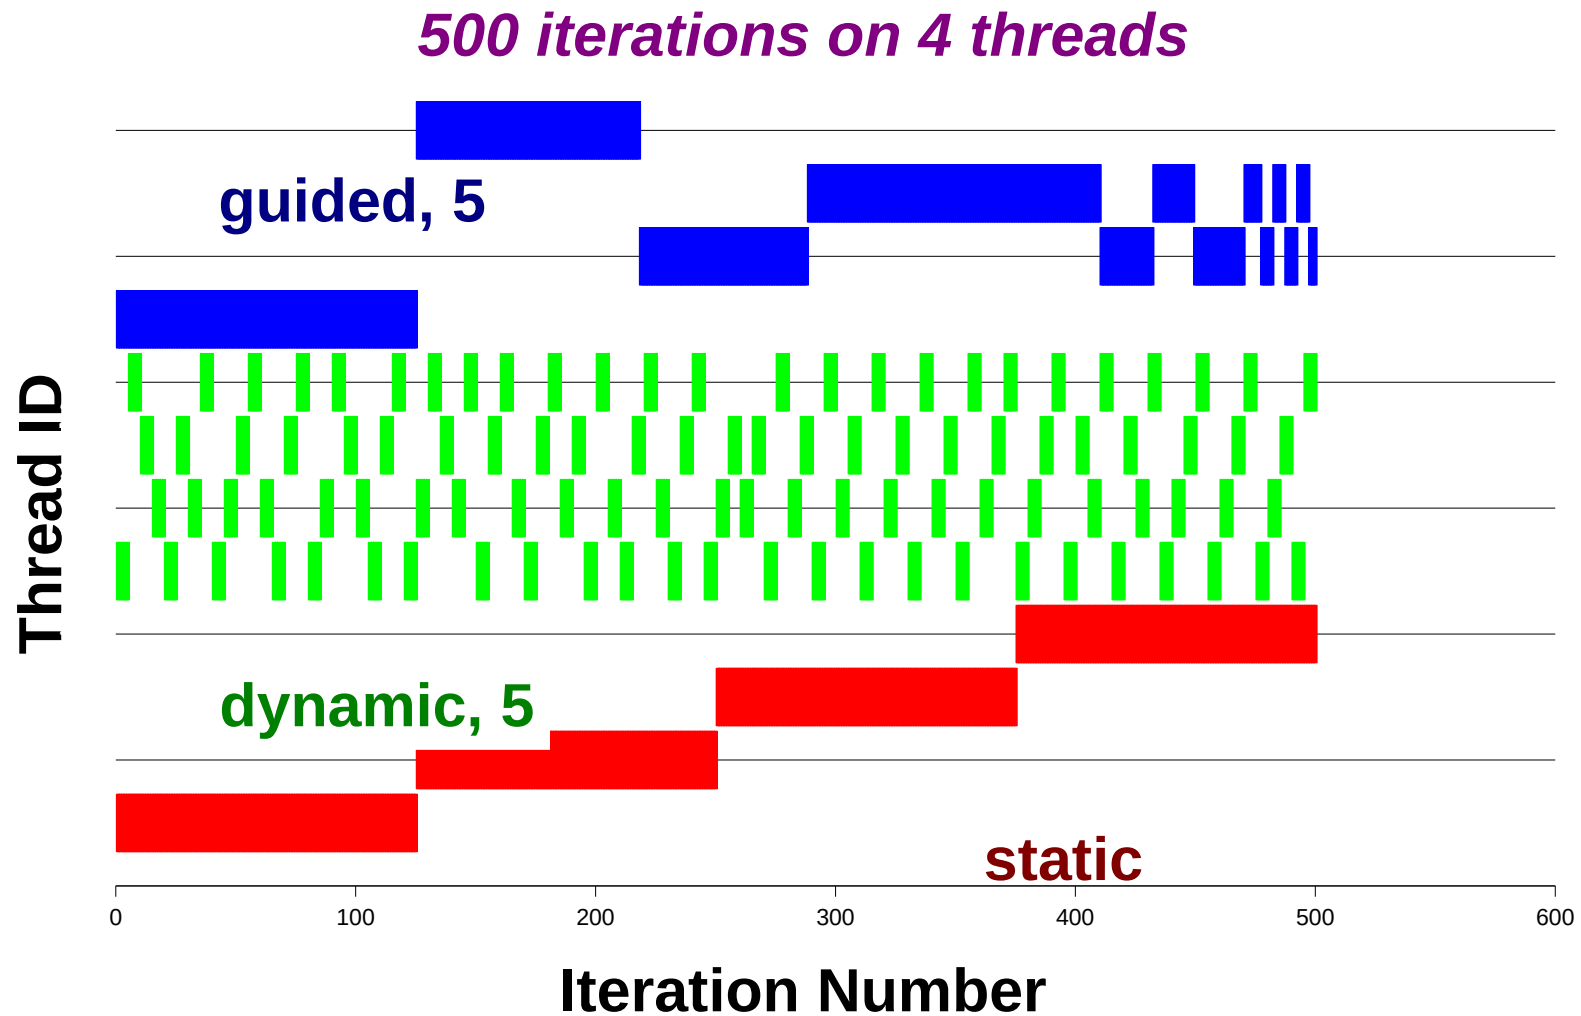
Guided schedule:

- ❑ The chunk size is exponentially reduced with each chunk that gets **dynamically** assigned to the threads; chunk defines the minimum number of iterations to assign each time.

$$\text{chunk} = \text{unass_iter} / (\text{weight} * \text{n_thr})$$

- ❑ The default minimum chunk size is 1.

OpenMP Scheduling



OpenMP Scheduling

```
#pragma omp for schedule(runtime)
```


Runtime schedule:

- ❑ The schedule is detected at runtime from the setting of the OMP_SCHEDULE environment variable.
- ❑ Syntax: OMP_SCHEDULE=type,chunk

A real world example: Molecular Dynamics simulation

Example: MD simulation

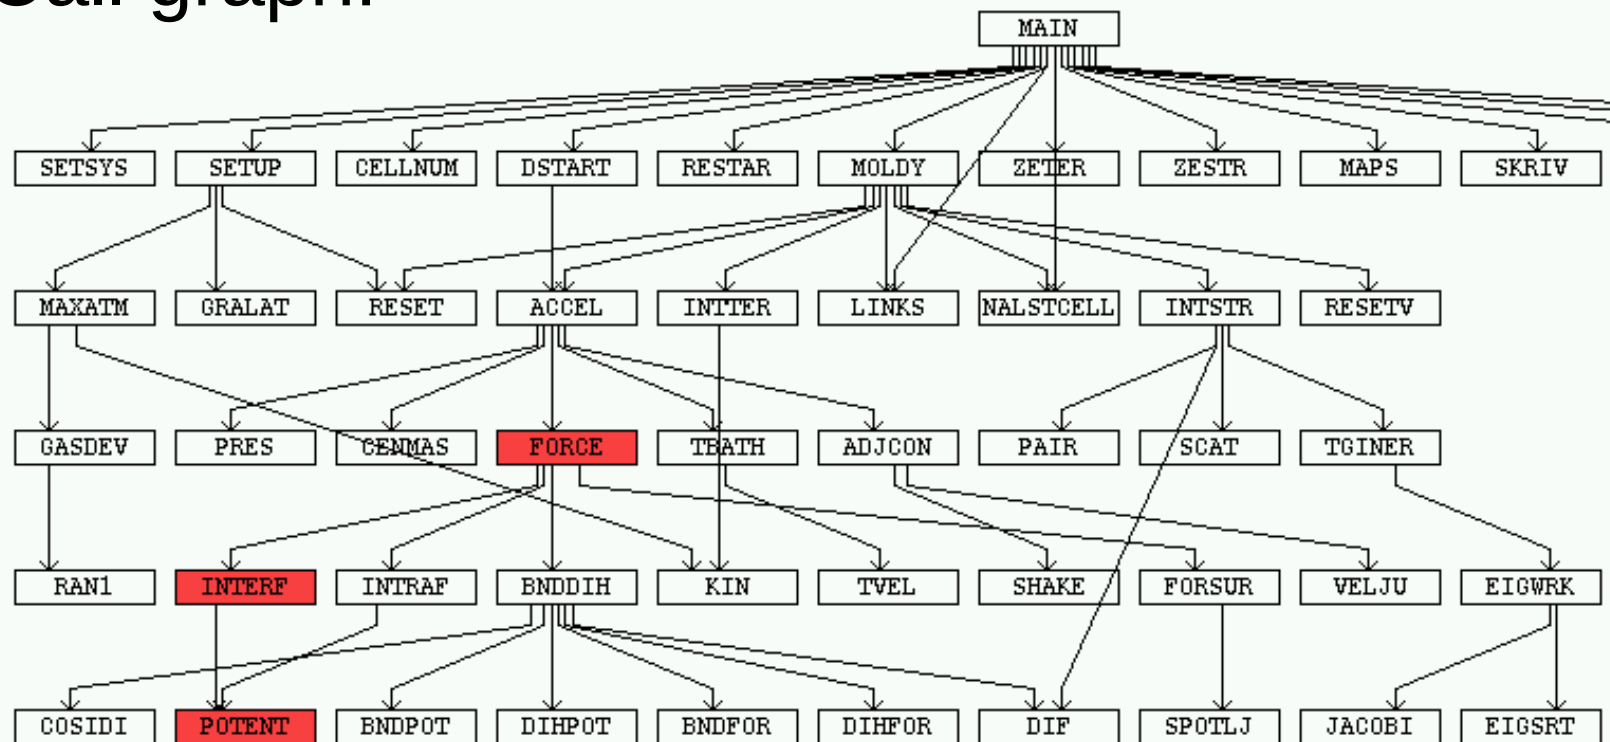
- ❑ Molecular Dynamics simulation of long carbon molecules on a surface:



- ❑ 7200+ lines of Fortran 77 code
- ❑ GOTOs, COMMON blocks, ...
- ❑ one source file

Example: MD simulation

Call graph:



more than 80% of the runtime are spent in the red part of the call graph

Example: MD simulation

- ❑ The loop to be parallelized contains a call to another subroutine.
- ❑ Data is passed the old Fortran style via COMMON blocks
- ❑ First try: Inserted one PARALLEL DO pragma in the code, using autoscoping, i.e. a feature of the Oracle Studio compiler
- ❑ The compiler generated a parallel version!

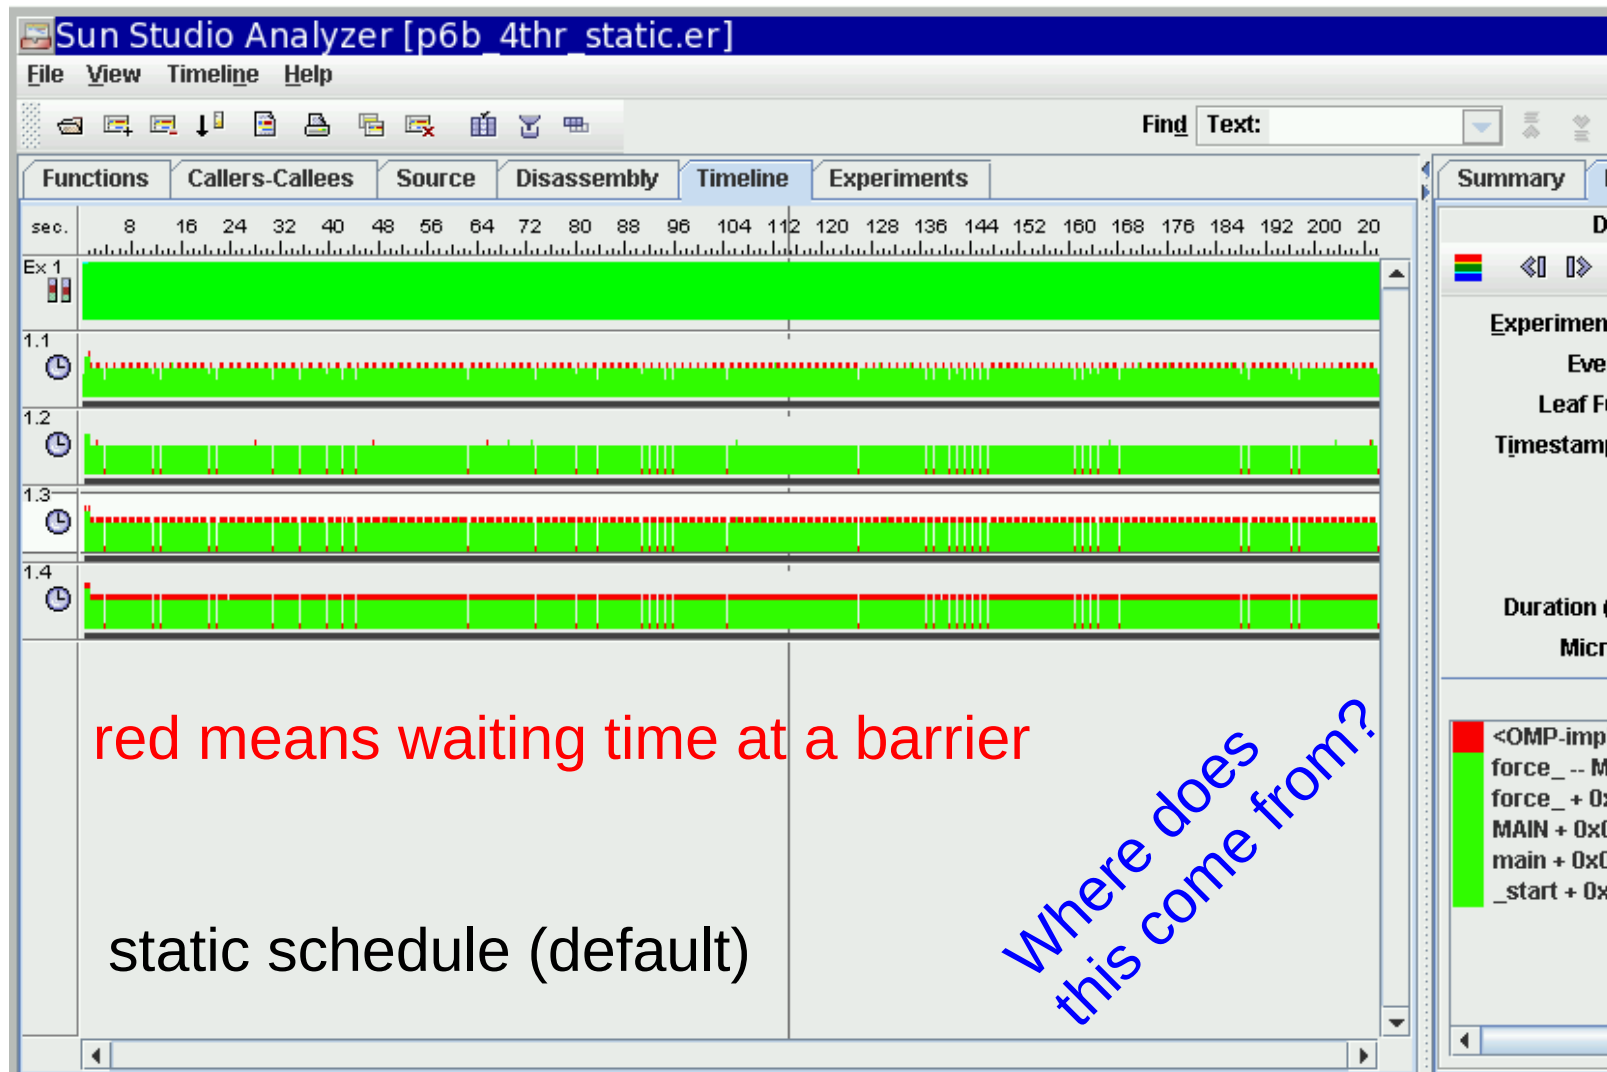
This took us by surprise!

Example: MD simulation

- ❑ First test runs:
 - ❑ It didn't scale ...
 - ❑ The results were dependent on the number of threads ...
- ❑ Thread analyzer revealed data races in two variables inside the called subroutine.
- ❑ Fix: Added additional scoping for those variables in the OpenMP pragma!
- ❑ This solved the data race problem.

Example: MD simulation


Analysis of the scaling problem:



Example: MD simulation

The simulation box:


seen from the side



subdivision into smaller cells

Example: MD simulation

Adapted the schedule:



Example: MD simulation

Speed-up results:

